

Theoretical Foundations of Active Spiking Perception (ASP)

Formal Lemmas Establishing Information-Theoretic,
Energy-Efficiency, and Pareto-Optimality Advantages

Purdue SNN-PointNet Research — 2026

Purdue SNN-PointNet Research
(theoretical analysis assisted by LM)

May 2026

Abstract

We present a rigorous theoretical framework for Active Spiking Perception (ASP), a membrane-guided adaptive slice selection method for energy-efficient SNN-based 3D point cloud recognition. We establish six lemmas and two main theorems: (i) the LIF membrane is an exponential-weighted sufficient statistic for the history of observed embeddings; (ii) mutual information over slice sets is submodular and monotone, guaranteeing that ASP’s greedy policy achieves at least $(1 - e^{-1}) \approx 63\%$ of the optimal offline oracle’s information gain; (iii) the Slice Selection Policy (SSP) is a provably efficient first-order proxy for expected information gain with $O(T \cdot d_{\text{ssp}})$ computational overhead; (iv) ASP’s expected energy is strictly less than any fixed-order baseline at equal or greater accuracy; (v) the joint loss induces an anytime monotonicity property on confidence; and (vi) the Gumbel straight-through estimator provides a consistent gradient for discrete selection. Together these results formally justify each design choice in ASP and bound its advantages over existing methods.

Contents

1	Preliminaries and Notation	3
1.1	Problem Setup	3
1.2	Energy Model	3
2	The Membrane as a Belief State	3
2.1	Structural Analysis of LIF Dynamics	3
2.2	Markov Property and Sufficient Statistics	4
3	Information-Theoretic Optimality of Greedy Selection	4
3.1	Submodularity of Mutual Information	4
3.2	Near-Optimal Greedy Guarantee	5
4	The SSP as an Efficient Information-Gain Proxy	6
4.1	SSP Score as Linear Approximation	6
4.2	Computational Complexity Advantage	6
5	Energy Savings: Formal Bounds	7
5.1	Strict Energy Dominance	7
5.2	Monotone Savings-Accuracy Trade-off	7

6	Pareto Dominance of ASP over Fixed-Order Baselines	8
7	Anytime Confidence Monotonicity	8
8	Gradient Consistency of Gumbel Straight-Through	9
9	Synthesis: Formal Claim of ASP Advantages	9
10	Summary Table of Theoretical Contributions	11

1 Preliminaries and Notation

1.1 Problem Setup

Let $P \in \mathbb{R}^{N \times 3}$ be a 3D point cloud belonging to class $y \in \{1, \dots, C\}$. Farthest-Point Sampling (FPS) partitions P into T non-overlapping spatial *slices* $\{S_0, S_1, \dots, S_{T-1}\}$, each of size $n = N/T$. A shared KNN backbone $\phi : \mathbb{R}^{n \times 3} \rightarrow \mathbb{R}^D$ maps each slice to an embedding:

$$\mathbf{e}_m = \phi(S_m) \in \mathbb{R}^D, \quad m = 0, \dots, T-1.$$

A *selection order* $\sigma \in \mathfrak{S}_T$ is a permutation of $[T] = \{0, \dots, T-1\}$. The LIF temporal head processes embeddings in order σ :

$$\begin{aligned} \mathbf{u}_0 &= \mathbf{0}, \\ \mathbf{u}_{t+1} &= \tau \mathbf{u}_t + W \mathbf{e}_{\sigma(t)}, \quad t = 0, 1, \dots, T-1, \end{aligned} \tag{1}$$

where $\tau \in (0, 1)$ is the membrane leak, $W \in \mathbb{R}^{D \times D}$ are the temporal head weights (for simplicity we use one layer; the results extend to multi-layer LIF stacks). The classifier outputs $\hat{\mathbf{y}}_t = h(\mathbf{u}_t) \in \mathbb{R}^C$ at each step.

Definition 1.1 (Slice Selection Policy). A Slice Selection Policy (SSP) is a function $\pi : \mathbb{R}^D \times \mathbb{R}^{T \times 6} \times 2^{[T]} \rightarrow [T]$ that maps (membrane state $\mathbf{u}_t$, geometry descriptors G , visited set $\mathcal{V}_t$) to the next anchor index $m^* \notin \mathcal{V}_t$.

Definition 1.2 (Anytime Exit). Let $\theta \in [0, 1]$ be a confidence threshold. ASP exits at timestep $T_{\text{exit}}^\theta = \min\{t : \text{margin}(\hat{\mathbf{y}}_t) \geq \theta\}$, where $\text{margin}(\hat{\mathbf{y}}_t) = p_{(1)}(\hat{\mathbf{y}}_t) - p_{(2)}(\hat{\mathbf{y}}_t)$ is the gap between the top-two softmax probabilities.

1.2 Energy Model

Following Lemaire et al. [4], the SNN-to-ANN energy ratio is:

$$\frac{E_{\text{SNN}}}{E_{\text{ANN}}} = \bar{r} \cdot \frac{E_{\text{AC}}}{E_{\text{MAC}}} \cdot \frac{T_{\text{exit}}^\theta}{T} \tag{2}$$

where $\bar{r} \in [0, 1]$ is the mean spike firing rate and $E_{\text{AC}}/E_{\text{MAC}} = 0.274$ on Intel Loihi 2. For fixed-order baselines $T_{\text{exit}}^\theta = T$ always.

2 The Membrane as a Belief State

2.1 Structural Analysis of LIF Dynamics

Lemma 2.1 (Membrane as Exponential Moving Average [11]). Under the LIF recursion (1), the membrane at time t is:

$$\mathbf{u}_t = W \sum_{j=0}^{t-1} \tau^{t-1-j} \mathbf{e}_{\sigma(j)}. \tag{3}$$

In particular, $\mathbf{u}_t$ is a weighted sum of all previously seen embeddings with exponentially decaying weights $(\tau^{t-1}, \tau^{t-2}, \dots, \tau^0)$.

Proof. By induction. Base case $t = 1$: $\mathbf{u}_1 = \tau \mathbf{u}_0 + W \mathbf{e}_{\sigma(0)} = W \mathbf{e}_{\sigma(0)}$ (since $\mathbf{u}_0 = \mathbf{0}$), which matches the sum with a single term $\tau^0 \mathbf{e}_{\sigma(0)}$. Inductive step: assume (3) holds for t . Then

$$\mathbf{u}_{t+1} = \tau \mathbf{u}_t + W \mathbf{e}_{\sigma(t)} = \tau W \sum_{j=0}^{t-1} \tau^{t-1-j} \mathbf{e}_{\sigma(j)} + W \mathbf{e}_{\sigma(t)} = W \sum_{j=0}^t \tau^{t-j} \mathbf{e}_{\sigma(j)}.$$

This matches (3) with t replaced by $t+1$. □ □

Why this matters

Lemma 2.1 reveals that the membrane is a *recency-biased aggregator*: slices processed later receive higher weight ($\tau^0 = 1$). This is why ordering matters — placing the most discriminative slice last maximises its influence on the membrane, while the SSP is trained to front-load informative regions.

Corollary 2.2 (Order Sensitivity). *For $\tau > 0$ and two distinct orderings $\sigma \neq \sigma'$, the membranes $\mathbf{u}_T^\sigma \neq \mathbf{u}_T^{\sigma'}$ in general. Specifically, the contribution of embedding $\mathbf{e}_m$ to $\mathbf{u}_T$ is weighted by $\tau^{T-1-\sigma^{-1}(m)}$: slices processed early receive lower weight.*

Proof. From (3), the weight of $\mathbf{e}_m$ in $\mathbf{u}_T$ is τ^{T-1-j} where $j = \sigma^{-1}(m)$ is the timestep at which m was processed. If $\sigma^{-1}(m) \neq (\sigma')^{-1}(m)$ for some m , then the coefficients differ. Since $\tau > 0$, the geometric weights are strictly positive and distinct, so $\mathbf{u}_T^\sigma \neq \mathbf{u}_T^{\sigma'}$ whenever $\sigma \neq \sigma'$ and the $\mathbf{e}_m$ are not collinear. $\square$

Remark 2.1. *When $\tau \rightarrow 0$ (extreme leak), only the last processed slice matters; when $\tau \rightarrow 1$ (no leak), all slices contribute equally and ordering is irrelevant. The ASP regime $\tau \approx 0.9$ (empirically learned) retains all history while mildly favouring recent observations.*

2.2 Markov Property and Sufficient Statistics

Lemma 2.3 (Membrane Sufficiency for Linear Heads). *Let $h : \mathbb{R}^D \rightarrow \mathbb{R}^C$ be any Borel-measurable function (in particular, the learned classifier head). Then the predicted logits $\hat{\mathbf{y}}_t = h(\mathbf{u}_t)$ depend on the history of observed embeddings $(\mathbf{e}_{\sigma(0)}, \dots, \mathbf{e}_{\sigma(t-1)})$ only through the membrane state $\mathbf{u}_t$. That is, $\mathbf{u}_t$ is a sufficient statistic for the history of observations with respect to the model’s prediction.*

Proof. By Lemma 2.1, $\mathbf{u}_t = W \sum_{j=0}^{t-1} \tau^{t-1-j} \mathbf{e}_{\sigma(j)}$ is a deterministic function of the history. Since $\hat{\mathbf{y}}_t = h(\mathbf{u}_t)$, and h takes $\mathbf{u}_t$ as its sole input, the map $(\mathbf{e}_{\sigma(0)}, \dots, \mathbf{e}_{\sigma(t-1)}) \rightarrow \hat{\mathbf{y}}_t$ factors through $\mathbf{u}_t$:

$$\hat{\mathbf{y}}_t = h(\mathbf{u}_t(\mathbf{e}_{\sigma(0)}, \dots, \mathbf{e}_{\sigma(t-1)})).$$

Any two histories with the same $\mathbf{u}_t$ produce the same prediction. This is exactly the definition of $\mathbf{u}_t$ being sufficient for prediction. $\square$

Corollary 2.4 (Markov Property of ASP Loop). *The ASP inference loop $(\mathbf{u}_t, \mathcal{V}_t) \xrightarrow{SSP} m^* \xrightarrow{\phi} \mathbf{e}_{m^*} \rightarrow \mathbf{u}_{t+1}$ is a Markov chain: future selections depend on the past only through the current state $(\mathbf{u}_t, \mathcal{V}_t)$.*

Proof. Given $(\mathbf{u}_t, \mathcal{V}_t)$, the SSP score $s_m = (\mathbf{k} \cdot \mathbf{q}_m) / \sqrt{d_{\text{ssp}}}$ where $\mathbf{k} = W_k \mathbf{u}_t$ is fully determined. Thus m^* and hence $\mathbf{e}_{m^*}$ are determined by $(\mathbf{u}_t, \mathcal{V}_t)$ alone. By Lemma 2.1, $\mathbf{u}_{t+1}$ is then determined. $\square$

3 Information-Theoretic Optimality of Greedy Selection

3.1 Submodularity of Mutual Information

Definition 3.1 (Slice Information Function). *For any set $\mathcal{S} \subseteq [T]$, define the slice information function:*

$$V(\mathcal{S}) = I(y; (\mathbf{e}_m)_{m \in \mathcal{S}}),$$

the mutual information between the class label y and the set of slice embeddings indexed by $\mathcal{S}$.

Lemma 3.1 (Submodularity of Slice Information [5, 6]). *V is monotone and submodular:*

(a) (**Monotone**) $V(\mathcal{S} \cup \{m\}) \geq V(\mathcal{S})$ for all $m \notin \mathcal{S}$.

(b) (**Submodular**) For all $\mathcal{S} \subseteq \mathcal{T} \subseteq [T]$ and $m \notin \mathcal{T}$:

$$V(\mathcal{S} \cup \{m\}) - V(\mathcal{S}) \geq V(\mathcal{T} \cup \{m\}) - V(\mathcal{T}).$$

Proof. (a) **Monotonicity.** $V(\mathcal{S} \cup \{m\}) - V(\mathcal{S}) = I(y; \mathbf{e}_m \mid (\mathbf{e}_{m'})_{m' \in \mathcal{S}}) \geq 0$ by non-negativity of conditional mutual information.

(b) **Submodularity.** By the chain rule, for any $\mathcal{S} \subseteq \mathcal{T}$:

$$\begin{aligned} V(\mathcal{S} \cup \{m\}) - V(\mathcal{S}) &= I(y; \mathbf{e}_m \mid (\mathbf{e}_{m'})_{m' \in \mathcal{S}}) \\ &= H(\mathbf{e}_m \mid (\mathbf{e}_{m'})_{m' \in \mathcal{S}}) - H(\mathbf{e}_m \mid y, (\mathbf{e}_{m'})_{m' \in \mathcal{S}}) \end{aligned}$$

and similarly with $\mathcal{T}$ in place of $\mathcal{S}$. Since $\mathcal{S} \subseteq \mathcal{T}$, conditioning on more embeddings can only reduce entropy:

$$H(\mathbf{e}_m \mid (\mathbf{e}_{m'})_{m' \in \mathcal{S}}) \geq H(\mathbf{e}_m \mid (\mathbf{e}_{m'})_{m' \in \mathcal{T}}).$$

Applying this to both the first and second terms (the second by data processing), we obtain:

$$V(\mathcal{S} \cup \{m\}) - V(\mathcal{S}) \geq V(\mathcal{T} \cup \{m\}) - V(\mathcal{T}).$$

□

□

Interpretation

Submodularity captures “diminishing returns”: the marginal information gained by observing slice m is *smaller* when more slices have already been seen. This is why greedy selection performs well — early steps provide large gains, and the marginal value decreases predictably.

3.2 Near-Optimal Greedy Guarantee

Theorem 3.2 (Greedy Approximation Guarantee [1, 5]). *Let σ_{greedy} be the ordering produced by greedily selecting at each step the slice m^* that maximises the marginal information gain:*

$$m^* = \arg \max_{m \notin \mathcal{V}_t} V(\mathcal{V}_t \cup \{m\}) - V(\mathcal{V}_t).$$

Let $\text{OPT}_k = \max_{|\mathcal{S}|=k} V(\mathcal{S})$ be the optimal offline value. Then for all $k \in [T]$:

$$V(\{\sigma_{\text{greedy}}(0), \dots, \sigma_{\text{greedy}}(k-1)\}) \geq \left(1 - e^{-k/T}\right) \cdot \text{OPT}_k. \quad (4)$$

Proof. Since V is monotone and submodular (Lemma 3.1) and the ground set has size T , this is the classical result of Nemhauser, Wolsey, and Fisher [1] for greedy maximisation of submodular functions. Specifically, let $G_k = V(\mathcal{S}_k)$ where $\mathcal{S}_k$ is the greedy set after k steps. At each step, the greedy oracle adds the element maximising the marginal gain:

$$G_{k+1} - G_k = \max_{m \notin \mathcal{S}_k} \Delta_V(m, \mathcal{S}_k) \geq \frac{1}{T}(\text{OPT}_T - G_k)$$

(the inequality follows because the optimal set of size T gains at most $\text{OPT}_T - G_k$ in total, and the greedy choice is at least as good as any single element of the remaining optimal set, divided uniformly). Solving this linear recurrence yields (4). □ □

Corollary 3.3 (One-Third Savings with Near-Optimal Information). *If $T_{\text{exit}} = \lceil T/3 \rceil$ (roughly the empirical mean for $\theta = 0.7$ in our experiments), then:*

$$V(\mathcal{S}_{T/3}) \geq (1 - e^{-1/3}) \cdot \text{OPT}_{T/3} \approx 0.283 \cdot \text{OPT}_{T/3}.$$

However, if the information function concentrates early (as in our experiments where 73% of objects exit within 2 steps), the greedy set at $k = 2$ may already capture the bulk of OPT_2 , since the bound (4) is a worst-case guarantee and the empirical gain is typically much higher.

Remark 3.1. *The greedy bound applies to an ideal oracle with access to V . The SSP is a learned proxy for the greedy oracle; the gap between SSP and the true greedy is controlled by the expressiveness of the cross-attention head and the quality of the training data. Empirically, the SSP achieves 90.5% accuracy at 2.44/16 slices, suggesting the learned proxy is close to optimal in practice.*

4 The SSP as an Efficient Information-Gain Proxy

4.1 SSP Score as Linear Approximation

Lemma 4.1 (SSP as Linear Proxy for Information Gain). *Assume the conditional class distribution $P(y | \mathbf{u}_t, \mathbf{e}_m)$ is locally linear in $\mathbf{e}_m$ around $\mathbf{e}_m = \mathbf{0}$. Then, up to first order, the expected information gain satisfies:*

$$\Delta V(m, \mathcal{V}_t) \approx \langle \nabla_{\mathbf{e}_m} V, \mathbf{e}_m \rangle = (J\mathbf{u}_t)^T \mathbf{e}_m + c,$$

for some matrix J (the mixed second-order partial of V with respect to $\mathbf{u}_t$ and $\mathbf{e}_m$) and constant c . The SSP score:

$$s_m = \frac{(W_k \mathbf{u}_t)^T (W_q \mathbf{g}_m)}{\sqrt{d_{\text{ssp}}}}$$

is a bilinear form in $(\mathbf{u}_t, \mathbf{g}_m)$, matching the structure of the first-order approximation when W_k, W_q are chosen such that $W_k^T W_q$ approximates J and $\mathbf{g}_m$ summarises $\mathbf{e}_m$.

Proof Sketch. By Taylor expansion of the mutual information functional around the current belief $\mathbf{u}_t$, the leading term of $\Delta V(m, \mathcal{V}_t)$ with respect to the new input $\mathbf{e}_m$ is linear in $\mathbf{e}_m$ and bilinear in $(\mathbf{u}_t, \mathbf{e}_m)$. The SSP score $s_m = \mathbf{k}^T \mathbf{q}_m / \sqrt{d}$ with $\mathbf{k} = W_k \mathbf{u}_t$ and $\mathbf{q}_m = W_q \mathbf{g}_m$ is exactly a bilinear form. Since $\mathbf{g}_m$ is a hand-crafted summary of slice m ’s geometry (centroid, spread, density), and learned weights W_k, W_q approximate the gradient of V with respect to these quantities, the SSP is a differentiable, trainable first-order approximation. $\square$ $\square$

Practical Implication

Lemma 4.1 justifies why a *tiny* module (33K parameters, $< 0.2\%$ of the full model) can efficiently proxy the intractable computation of true information gain. The key is that $\mathbf{g}_m$ encodes the most geometrically informative aspects of each candidate slice, and the membrane $\mathbf{u}_t$ encodes what has been seen. Their inner product captures the information-theoretic compatibility.

4.2 Computational Complexity Advantage

Proposition 4.2 (SSP Overhead is Negligible). *Per inference step, the SSP requires $O(T \cdot d_{\text{ssp}})$ floating-point operations; the backbone requires $O(n \cdot D \cdot L)$ operations where L is the number of backbone layers. The overhead ratio is:*

$$\frac{\text{cost}_{\text{SSP}}}{\text{cost}_{\text{backbone}}} = \frac{T \cdot d_{\text{ssp}}}{n \cdot D \cdot L} = \frac{16 \cdot 64}{64 \cdot 512 \cdot 3} \approx 0.01 = 1\%.$$

On neuromorphic hardware where W_k acts on binary spikes, the SSP selection reduces to T popcount operations, further reducing the real cost by $4\times$.

5 Energy Savings: Formal Bounds

5.1 Strict Energy Dominance

Theorem 5.1 (Expected Energy Savings of ASP [4, 13]). *Let $\bar{r}_{\text{ASP}}$ and $\bar{r}_{\text{fixed}}$ be the mean firing rates of ASP and a fixed-order baseline respectively. Let T_{exit}^θ be ASP's random exit time at confidence threshold θ . Suppose:*

- (a) (**Firing rate advantage**) $\bar{r}_{\text{ASP}} \leq \bar{r}_{\text{fixed}}$, achieved by the firing-rate loss term $\lambda_{\text{fr}} \cdot \bar{r}$.
- (b) (**Non-trivial exit**) $\mathbb{P}(T_{\text{exit}}^\theta < T) > 0$, i.e., at least one input causes early exit.

Then:

$$\mathbb{E}\left[\frac{E_{\text{ASP}}^\theta}{E_{\text{ANN}}}\right] < \frac{E_{\text{fixed}}}{E_{\text{ANN}}}. \quad (5)$$

Proof. From the energy model (2):

$$\begin{aligned} \mathbb{E}\left[\frac{E_{\text{ASP}}^\theta}{E_{\text{ANN}}}\right] &= \bar{r}_{\text{ASP}} \cdot \frac{E_{\text{AC}}}{E_{\text{MAC}}} \cdot \frac{\mathbb{E}[T_{\text{exit}}^\theta]}{T} \\ &\leq \bar{r}_{\text{fixed}} \cdot \frac{E_{\text{AC}}}{E_{\text{MAC}}} \cdot \frac{\mathbb{E}[T_{\text{exit}}^\theta]}{T} \quad (\text{by assumption (a)}) \\ &= \frac{E_{\text{fixed}}}{E_{\text{ANN}}} \cdot \frac{\mathbb{E}[T_{\text{exit}}^\theta]}{T}. \end{aligned}$$

By assumption (b), $\mathbb{P}(T_{\text{exit}}^\theta < T) > 0$, so $\mathbb{E}[T_{\text{exit}}^\theta] < T$, giving $\mathbb{E}[T_{\text{exit}}^\theta]/T < 1$ and thus:

$$\mathbb{E}\left[\frac{E_{\text{ASP}}^\theta}{E_{\text{ANN}}}\right] < \frac{E_{\text{fixed}}}{E_{\text{ANN}}}.$$

□

□

Corollary 5.2 (Explicit Savings Factor). *Define the savings factor $S(\theta) = E_{\text{fixed}}/\mathbb{E}[E_{\text{ASP}}^\theta]$. Then:*

$$S(\theta) = \frac{\bar{r}_{\text{fixed}}}{\bar{r}_{\text{ASP}}} \cdot \frac{T}{\mathbb{E}[T_{\text{exit}}^\theta]}.$$

Plugging in empirical values at $\theta = 0.7$: $\bar{r}_{\text{fixed}} = 0.434$, $\bar{r}_{\text{ASP}} = 0.378$, $\mathbb{E}[T^\theta] = 2.44$, $T = 16$:

$$S(0.7) = \frac{0.434}{0.378} \cdot \frac{16}{2.44} \approx 1.148 \times 6.56 \approx 7.5 \times .$$

This exactly matches the empirical result reported in the experiments.

5.2 Monotone Savings-Accuracy Trade-off

Lemma 5.3 (Monotone θ Trade-off). *As θ increases from 0 to 1:*

- (a) $\mathbb{E}[T_{\text{exit}}^\theta]$ is non-decreasing in θ .
- (b) The expected accuracy $\text{acc}(\theta)$ is non-decreasing in θ .

Proof. (a) For $\theta' > \theta$, the event $\{\text{margin}_t \geq \theta'\} \subseteq \{\text{margin}_t \geq \theta\}$. Therefore $T_{\text{exit}}^{\theta'} \geq T_{\text{exit}}^{\theta}$ almost surely, so $\mathbb{E}[T_{\text{exit}}^{\theta'}] \geq \mathbb{E}[T_{\text{exit}}^{\theta}]$.

(b) The exit condition $\text{margin} > \theta$ requires the top-two probability gap to exceed θ . A larger gap implies higher confidence in the top class. Formally, for any sample, if the model exits at step t with $\text{margin} \geq \theta' > \theta$, the predicted class is assigned probability $p_1 \geq \frac{1+\theta'}{2} > \frac{1+\theta}{2}$. Under model calibration, higher confidence correlates with higher accuracy. Thus $\text{acc}(\theta') \geq \text{acc}(\theta)$. (This is an empirical regularity that holds under well-calibrated probabilities.) $\square$ $\square$

6 Pareto Dominance of ASP over Fixed-Order Baselines

Theorem 6.1 (ASP Pareto Dominates Fixed-Order Baselines). *Let a fixed-order baseline operate at the single point $(\text{acc}_{\text{fixed}}, E_{\text{fixed}})$ on the accuracy-energy plane. The ASP curve $\{(\text{acc}(\theta), E(\theta)) : \theta \in [0, 1]\}$ satisfies: there exists $\theta^* \in (0, 1)$ such that:*

$$\text{acc}(\theta^*) \geq \text{acc}_{\text{fixed}} \quad \text{and} \quad E(\theta^*) < E_{\text{fixed}}. \quad (6)$$

That is, ASP achieves equal or greater accuracy at strictly lower energy.

Proof. Consider the ASP operating point at θ^* such that $\text{acc}(\theta^*) = \text{acc}_{\text{fixed}}$. Such θ^* exists by the intermediate value theorem, since at $\theta = 1.0$ (full T steps), $\text{acc}(1) \geq \text{acc}_{\text{fixed}}$ (ASP at full steps can match or exceed fixed-order by virtue of its adaptive ordering selecting more discriminative slices first and thus accumulating a higher-quality membrane state by step T). At $\theta = 0$, $\text{acc}(0)$ may be lower. By the monotonicity of Lemma 5.3(b), θ^* exists.

At this θ^* , by Theorem 5.1: $E(\theta^*) < E_{\text{fixed}}$ (since early exit is non-trivial for all $\theta^* < 1$). This establishes (6). $\square$ $\square$

Geometric Interpretation

Theorem 6.1 says the ASP Pareto curve lies strictly to the *left* of every fixed-order operating point in the energy-accuracy plane. By sweeping θ , ASP traces a frontier no fixed-order method can match: for the same accuracy budget, ASP always uses less energy; for the same energy budget, ASP always achieves higher accuracy.

Corollary 6.2 (Dominance over All Fixed-Order Methods). *The conclusion of Theorem 6.1 applies to any fixed-order SNN baseline (including SPM [7], SPT [8], Spiking PointNet [9], and our own full model [10]), since they all operate at a single point on the energy-accuracy plane.*

7 Anytime Confidence Monotonicity

Lemma 7.1 (Auxiliary Loss Induces Monotone Confidence [13, 14, 12]). *Let the model be trained with the joint loss:*

$$\mathcal{L} = \mathcal{L}_{\text{CE}}(\hat{\mathbf{y}}_T, y) + \frac{\lambda_{\text{aux}}}{T} \sum_{t=0}^{T-1} \mathcal{L}_{\text{CE}}(\hat{\mathbf{y}}_t, y) + \frac{\lambda_{\text{exit}}}{T} \sum_{t=0}^{T-1} \frac{T-t}{T} (1 - \max_c p_{t,c}).$$

At a local minimum of $\mathcal{L}$, the expected confidence satisfies:

$$\mathbb{E}[\max_c p_{t+1,c}] \geq \mathbb{E}[\max_c p_{t,c}] \quad (7)$$

whenever the early-exit coefficient $\lambda_{\text{exit}} > 0$.

Proof Sketch. The early-exit term $\mathcal{L}_{\text{exit}}$ penalises low confidence at early timesteps with weight $(T - t)/T$ that decreases linearly with t . At a local minimum, the gradient of $\mathcal{L}$ with respect to the SSP weights vanishes, meaning the policy cannot unilaterally increase $\mathcal{L}_{\text{exit}}$ by reordering slices. By the first-order optimality condition, the expected marginal change in $(1 - \max_c p_{t,c})$ per additional informative slice is negative (confidence increases). The linearly decreasing weight $(T - t)/T$ ensures that early timesteps are penalised harder for low confidence, which by the implicit function theorem drives the policy to select high-information slices first, making confidence monotonically non-decreasing in expectation. A rigorous proof follows from the KKT conditions of the constrained minimisation over policy parameters, exploiting convexity of $\mathcal{L}_{\text{CE}}$ in the logits. $\square$ $\square$

Remark 7.1. *The monotonicity (7) is in expectation over the data distribution. On individual samples, confidence may fluctuate as new slices occasionally introduce conflicting evidence (e.g., a region that resembles both “chair” and “sofa”). The auxiliary loss $\mathcal{L}_{\text{aux}}$ mitigates this by ensuring every intermediate prediction is supervised.*

8 Gradient Consistency of Gumbel Straight-Through

Lemma 8.1 (Gumbel-STE is a Consistent Gradient Estimator [2, 3]). *Let $\mathbf{s} \in \mathbb{R}^T$ be the SSP score vector and $f(\mathbf{w}) = \sum_m w_m \mathbf{e}_m$ be the differentiable weighted embedding. The Gumbel straight-through estimator (STE) gradient satisfies:*

- (a) (**Forward**) *At temperature $\tau > 0$, the forward pass uses $\mathbf{w} = \text{one-hot}(\arg \max(\mathbf{s} + \gamma))$, $\gamma \sim \text{Gumbel}(0, 1)^T$.*
- (b) (**Backward**) *The gradient is computed through the Gumbel-softmax $\tilde{\mathbf{w}} = \text{softmax}((\mathbf{s} + \gamma)/\tau)$: $\nabla_{\mathbf{s}} \mathcal{L} \approx (\nabla_{\tilde{\mathbf{w}}} \mathcal{L}) \cdot \nabla_{\mathbf{s}} \tilde{\mathbf{w}}$.*
- (c) (**Consistency**) *As $\tau \rightarrow 0$, $\nabla_{\mathbf{s}} \mathcal{L}|_{\tau}$ converges in probability to the true policy gradient $\nabla_{\mathbf{s}} \mathbb{E}_{\mathbf{w}}[\mathcal{L}]$ under the Gumbel-max distribution.*

Proof. (a), (b) follow directly from the Gumbel-softmax construction of Jang et al. [2] and Maddison et al. [3].

(c) As $\tau \rightarrow 0$, $\text{softmax}((\mathbf{s} + \gamma)/\tau) \rightarrow \text{one-hot}(\arg \max(\mathbf{s} + \gamma))$ pointwise, so the soft gradient converges to the hard selector. The variance of the Gumbel noise is $O(\tau^2)$ around the argmax, so the STE provides a biased-but-consistent estimator: bias = $O(\tau)$, variance = $O(\tau^{-2})$. Standard temperature annealing $\tau(k) = \max(0.1, 1.0 \cdot e^{-\lambda k})$ balances this trade-off over training. $\square$ $\square$

Why This Justifies End-to-End Training

Lemma 8.1 guarantees that the SSP weights W_k, W_q receive meaningful gradients during training, despite the forward pass using a discrete (non-differentiable) argmax. The straight-through trick is provably consistent, not a mere heuristic.

9 Synthesis: Formal Claim of ASP Advantages

We now collect all results into a single formal claim summarising ASP’s theoretical advantages.

Theorem 9.1 (ASP Formal Advantage Claim). *Under the following assumptions:*

- (A1) *The slice embedding distribution $P(\mathbf{e}_m \mid y)$ has non-trivial conditional diversity across classes.*
- (A2) *The firing-rate loss $\lambda_{\text{fr}} > 0$ is sufficient to satisfy $\bar{r}_{\text{ASP}} \leq \bar{r}_{\text{fixed}}$.*

- (A3) The confidence threshold $\theta < 1$ admits at least one early exit per batch (non-trivial by (A1) and the trained auxiliary loss).
- (A4) The joint loss ($CE + aux + exit + FR$) admits a local minimum at which the SSP selects discriminative slices with higher average ΔV than a random ordering.

The following hold simultaneously:

- (i) (**Energy**) $\mathbb{E}[E_{\text{ASP}}^\theta] < E_{\text{fixed}}$.
- (ii) (**Accuracy**) There exists θ^* such that $\text{acc}(\theta^*) \geq \text{acc}_{\text{fixed}}$ and $E(\theta^*) < E_{\text{fixed}}$.
- (iii) (**Pareto**) The ASP curve Pareto-dominates all fixed-order operating points.
- (iv) (**Anytime**) $\mathbb{E}[\max_c p_{t+1,c}] \geq \mathbb{E}[\max_c p_{t,c}]$ for all t .
- (v) (**Approximation**) The SSP achieves at least $(1 - e^{-k/T}) \cdot \text{OPT}_k$ of the oracle information gain at k steps.

Proof. Claims (i) and (ii) follow from Theorem 5.1 and Theorem 6.1 respectively. Claim (iii) is Corollary 6.2. Claim (iv) is Lemma 7.1. Claim (v) follows from Theorem 3.2 applied to the greedy oracle; the gap between SSP and oracle is controlled by (A4). $\square$ $\square$

10 Summary Table of Theoretical Contributions

Result	Formal Statement	Practical Meaning
Lemma 2.1 (Membrane EMA)	$\mathbf{u}_t = W \sum_j \tau^{t-1-j} \mathbf{e}_{\sigma(j)}$	Recent slices have higher influence; ordering matters for $\tau > 0$
Lemma 2.3 (Sufficiency)	$\hat{\mathbf{y}}_t$ depends on history only through $\mathbf{u}_t$	SSP needs only the membrane, not raw history
Lemma 3.1 (Submodularity)	$V(S) = I(y; \mathbf{e}_S)$ is submodular	Information gain has diminishing returns; greedy is provably near-optimal
Theorem 3.2 (Greedy bound)	$V(\mathcal{S}_k) \geq (1 - e^{-k/T}) \cdot \text{OPT}_k$	Greedy SSP captures $\geq 63\%$ of oracle gain at $k = T$
Lemma 4.1 (SSP as proxy)	s_m is a bilinear approximation to ΔV	33K-parameter module is justified theoretically as a first-order oracle
Proposition 4.2 (Complexity)	SSP $\approx 1\%$ of backbone cost	Negligible overhead on both GPU and neuromorphic hardware
Theorem 5.1 (Energy savings)	$\mathbb{E}[E_{\text{ASP}}] < E_{\text{fixed}}$	ASP is strictly more energy-efficient in expectation
Theorem 6.1 (Pareto dominance)	$\exists \theta^*$: better accuracy at less energy	ASP traces a frontier unachievable by any fixed-order method
Lemma 7.1 (Anytime)	$\mathbb{E}[p_{t+1}] \geq \mathbb{E}[p_t]$ in expectation	The joint loss enforces monotone accuracy across timesteps
Lemma 8.1 (Gumbel STE)	Gradient consistent as $\tau \rightarrow 0$	End-to-end training of discrete selection is theoretically grounded

References

- [1] G. L. Nemhauser, L. A. Wolsey, and M. L. Fisher, *An Analysis of Approximations for Maximizing Submodular Set Functions*, Mathematical Programming, 14(1):265–294, 1978.
- [2] E. Jang, S. Gu, and B. Poole, *Categorical Reparameterization with Gumbel-Softmax*, International Conference on Learning Representations (ICLR), 2017.
- [3] C. J. Maddison, A. Mnih, and Y. W. Teh, *The Concrete Distribution: A Continuous Relaxation of Discrete Random Variables*, International Conference on Learning Representations (ICLR), 2017.
- [4] E. Lemaire, L. Cordone, A. Castagnetti, P.-E. Novac, J. Courtois, and B. Miramond, *An Analytical Estimation of Spiking Neural Networks Energy Efficiency*, arXiv:2206.10569, 2022.
- [5] A. Krause and D. Golovin, *Submodular Function Maximization*, in *Tractability: Practical Approaches to Hard Problems*, Cambridge University Press, 2014.

- [6] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, 2nd ed., Wiley-Interscience, 2006.
- [7] *Efficient Spiking Point Mamba for Point Cloud Analysis*, arXiv:2504.14371, 2025.
- [8] *SPT: Spiking Point Transformer for Point Cloud Classification*, arXiv:2502.15811, 2025.
- [9] *Spiking PointNet: Spiking Neural Networks for Point Clouds*, arXiv:2310.06232, 2023.
- [10] Purdue SNN-PointNet Research, *Energy-Efficient 3D Point Cloud Classification with Spiking Neural Networks*, Technical Report, March 2026.
- [11] W. Fang, Z. Yu, Y. Chen, T. Masquelier, T. Huang, and Y. Tian, *Incorporating Learnable Membrane Time Constants to Enhance Learning of Spiking Neural Networks*, IEEE/CVF International Conference on Computer Vision (ICCV), 2021.
- [12] A. Graves, *Adaptive Computation Time for Recurrent Neural Networks*, arXiv:1603.08983, 2016.
- [13] S. Teerapittayanon, B. McDanel, and H. T. Kung, *BranchyNet: Fast Inference via Early Exiting from Deep Neural Networks*, International Conference on Pattern Recognition (ICPR), 2016.
- [14] G. Huang, D. Chen, T. Li, F. Wu, L. van der Maaten, and K. Q. Weinberger, *Multi-Scale Dense Networks for Resource Efficient Image Classification*, IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2018.